

Posto: Probabilistic Safety Verification Tool for Complex Autonomous Systems

Prachi Bhattacharjee^a, Bineet Ghosh^{✉a}, Étienne André^{✉b,c}

^a*University of Alabama, AL, USA*

^b*Nantes Université, Nantes, France*

^c*Institut universitaire de France (IUF)*

Abstract

The rapid advancement of autonomy in modern systems has led to a surge in model complexity, making safety verification using traditional tools increasingly difficult. In response, we introduce **Posto**, a statistical monitoring tool that performs safety verification of complex autonomous systems, including those driven by deep neural networks (DNN). While traditional monitoring tools depend on formal models to assess safety, **Posto** only requires the system’s input–output behavior (I/O execution) and the log (which may be noisy and incomplete). We provide easy-to-use scripts and comprehensive documentation to enable users to reproduce our experiments, monitor their own systems using logs, and extend the tool for their own application. We evaluate **Posto** through multiple case studies involving complex autonomous systems with DNN components.

Keywords: monitoring, neural networks, safety, monitoring tool, statistical verification

1. Motivation and significance

Over the past decade, distributed autonomous systems have been rapidly adopted across a wide range of domains. Although traditional formal verification tools have shown success in many safety-critical settings [3], they are increasingly struggling to cope with the rising complexity of these systems. In particular, these tools typically require a formal model—something that is often unavailable due to black-box components or the absence of a complete system model. To address these challenges, we introduce **Posto**¹ (see details in [Table 1](#)), a probabilistic monitoring tool that implements a statistical

¹Tool webpage: <https://autmn-lab.github.io/Posto/>

C1	Current code version	v1.0
C2	Permanent link to code/repository used for this code version	https://github.com/autmn-lab/posto_tool/releases/tag/v1.0
C3	Permanent link to Reproducible Capsule	https://doi.org/10.5281/zenodo.18233568
C4	Legal Code License	<i>GNU General Public License v3.0 (gpl-3.0)</i> [1]
C5	Code versioning system used	git
C6	Software code languages, tools, and services used	<code>Python, numpy, scipy, mpmath</code>
C7	Compilation requirements, operating environments and dependencies	See Posto Manual [2]
C8	Link to developer documentation/manual	See Posto Manual [2]
C9	Support email for questions	autmnlab.cs@gmail.com

Table 1: Code metadata

1 hypothesis testing based algorithm (see [4]) to assess system safety directly
 2 from the execution logs, even when the logs contain noisy or missing samples.
 3 Using only the system’s input–output behavior (independent of internal im-
 4 plementations, including DNN-based components), **Posto** can either certify
 5 the system as safe with a high probabilistic guarantee, or generate a concrete
 6 counterexample when a violation is detected.

7 *Contributions.* Monitoring has emerged as a lightweight yet practically effec-
 8 tive verification technique capable of handling systems of considerable com-
 9 plexity (e.g., [5, 6, 7, 8]). Statistical monitoring, in particular, aims to ensure
 10 that autonomous systems are trustworthy, which means that they are safe
 11 for most practical purposes with a high level of confidence. By relaxing strict
 12 formal safety guarantees in favor of statistical ones (using the method pro-
 13 posed in [4]), **Posto** provides two key advantages: *i*) it is broadly applicable
 14 and can accommodate a wide range of system types—including nonlinear
 15 and DNN-based models—by relying solely on the system’s input/output ex-
 16 ecution behavior; and *ii*) it allows users to specify the desired confidence
 17 level for the analysis. For e.g., in a safety critical scenario, a high confidence
 18 level (e.g., 0.9999) may be chosen, while in less critical situations, a lower
 19 confidence level can be selected. **Posto** requires less time for monitoring as
 20 the desired confidence level decreases.

21 *Related Work.* Signal Temporal Logic (STL) formalizes real-valued signal
 22 properties in dense-time contexts. In addition to STL-based monitoring, re-
 23 cent attention has shifted towards monitoring using automata-based spec-
 24 ifications. Techniques such as timed pattern matching on timed regular
 25 expressions [9, 10, 11] have emerged, particularly in the context of deter-

ministic offline monitoring. The concept of *model-bounded monitoring*, introduced in [12], deviates from the conventional approach of monitoring a black-box system solely against a specification. Instead, it incorporates a limited, over-approximated understanding of the system to mitigate false positives. The over-approximated knowledge takes the form of a *linear hybrid automaton* [13]. Model-bounded monitoring was extended in the framework of uncertain linear systems in [8]. Unlike all these deterministic approaches, our proposed monitoring approach does not require any formal model—and can handle complex models such as non-linear models, DNNs etc.—of the system. It trades off formal safety guarantees with high confidence probabilistic guarantees.

In addition, *statistical verification* has found extensive application across various domains [14].

Outline. We briefly recall the software in Section 2: we provide illustrative examples in Section 3; we discuss the potential impact of **Posto** in Section 4; we conclude in Section 5.

2. Software description

Posto is an open-source Python-based software designed to run on Linux platforms.

2.1. Software architecture

The architecture (and dataflow) is given in Fig. A.1. Each block corresponds to a core component of the tool’s architecture, with its specific role indicated in the diagram. We now wish to discuss the core architectural components of **Posto**, followed by a discussion of its key software functionalities and the role of these components in supporting them.

Inputs and Outputs. The main inputs to our tool are: the I/O execution system model, the log (with noisy and missing samples), and the safety specifications against which the system is to be analyzed (with a high probabilistic guarantee). Analyzing these inputs (in Block 5), the tool produces either a safety certificate, indicating that the system satisfies the safety conditions with a high probabilistic guarantee, or a counterexample that reconstructs a concrete system behavior of the system (from the given log) that violates the safety condition. In addition, it also generates plots illustrating both safe and unsafe behaviors that provides users with deeper insight into the system’s execution.

1 *Check Safety (**Block 5**)*. This is the main block of the tool that is responsible
2 to perform the safety verification of the system. It takes the safety constraints specified by the user as an input, along with the computed value
3 K from **Block 3**, which denotes the number of valid system trajectories that must be checked to achieve the required probabilistic guarantee. It is worth
4 noting that these trajectories are valid reconstructions of the system behavior from the given log, where the log is either generated synthetically (in
5 **Block 2**) or taken from a real system execution. Once the inputs are collected,
6 the block verifies whether *all* K valid trajectories satisfy the safety constraints. If they do, the tool asserts the system as safe with the required
7 probabilistic confidence. Otherwise, the trajectory violating the safety constraint is returned as a counterexample, which serves as a concrete system
8 behavior recreated from the log that violates safety.

14 *2.2. Software functionalities*

15 *Safety Verification*. **Posto** enables scalable safety verification of complex autonomous systems, including systems with DNN components, using logs that may contain missing and noisy samples. The verification is performed against user-defined safety specifications with high probabilistic guarantees. The tool takes the system’s I/O execution model and the log as inputs. **Posto** performs safety analysis by reconstructing random system trajectories from the given log. If no trajectory violating the safety specification is identified, the system is asserted to be safe with the required probabilistic guarantee. Otherwise, the trajectory violating safety is returned as a concrete (i.e., reconstructed from the log) counterexample.

25 *Log Generation*. While **Posto** is primarily designed for safety verification of a system using a user given log, it also supports the generation of synthetic logs that can be used for safety analysis. In addition, the tool provides visualizations of the generated logs, offering a more intuitive and physically meaningful understanding of the system’s execution in real-world conditions.

30 *Behavior Generation*. **Posto** not only supports safety verification of complex autonomous systems, but also provides diagnostic tools to analyze system behavior and evaluate its robustness to different logging parameters, making it useful for system designers.

34 **3. Illustrative examples**

35 In this section, we reproduce the experimental results reported in [4, Section 6] for the Jet model, Van der Pol oscillator and the Mountain Car case studies. These experiments are designed to assess the scalability and

1 practical usability of the tool on complex autonomous systems, including
2 systems with DNN-based components. In particular, the evaluation studies:
3 *i*) the impact of varying the logging probability (i.e., the number of records
4 in the log), *ii*) the impact of varying the amount of uncertainties in the log
5 samples, *iii*) impact of probabilistic confidence c , *iv*) effect of various other
6 parameters.

7 *Note on stochasticity of the recreations.* It is worth noting that the safety
8 verification method proposed in our work is inherently statistical. Conse-
9 quently, the reproduction of the plots will not yield an exact match but
10 rather a stochastic recreation. In other words, the plots are recreated using
11 the same set of parameters as in the draft, which define the distributions.
12 The recreated plots represent one possible outcome from that distribution,
13 hence the term “stochastic recreation”. As a result, some figures that were
14 initially inferred to be safe in the draft may yield unsafe results (or vice versa)
15 during this process.

16 *Jet Model.* The main results of the Jet Model study are provided in Figs. A.2
17 and A.3 (as in [4, Figs. 2 and 3]). Once the tool is downloaded and set up,
18 the experimental results can be reproduced using the script² provided in our
19 GitHub repository, with detailed usage instructions available in [15] (and
20 additional steps in Appendix A). For instance, to recreate the result shown
21 in Fig. A.2a (similarly, Figs. A.2b and A.2c, ..., Figs. A.3c and A.3d), one
22 can execute the following command:

23 `python artEval.py --fig=A2a`

24 *Van der Pol Oscillator.* The main results of the Van der Pol Oscillator study
25 are provided in Fig. A.4 (as in [4, Fig. 4]). The results for the Van der Pol
26 Oscillator can be reproduced in a manner similar to the Jet Case Study
27 (see Appendix A for details). For example, to recreate the result shown in
28 Fig. A.4a, one can run the following command:

29 `python artEval.py --fig=A4a`

30 *Mountain Car.* We also present additional experiments (beyond those re-
31 ported in [4]) using a DNN-based controller for the mountain car bench-
32 mark [16] in Appendix B. These experiments further demonstrate the efficacy
33 and applicability of the proposed tool.

²https://github.com/autmn-lab/posto_tool/blob/master/artEval.py

1 **4. Impact**

2 Although **Posto** is still a prototype tool based on a recent algorithm [4]—
3 and not yet an industry-adopted software—we believe it has potential to
4 develop a user base interested in monitoring complex autonomous systems
5 (including those with DNN-based and black-box components) against safety
6 properties.

7 **Posto** is capable of efficiently monitoring complex systems against safety
8 properties very efficiently. It operates directly on noisy logs (even with miss-
9 ing samples) while relying only on the I/O execution model of the system
10 to assess its safety. This makes **Posto** particularly suitable for modern au-
11 tonomous systems, where perfect logs and formal models are rarely available.
12 Given its broad applicability, minimal modeling requirements, and computa-
13 tional efficiency, we believe **Posto** has the potential to open up new research
14 directions in the verification of complex autonomous systems, especially in
15 the context of statistical and data-driven safety analysis.

16 Furthermore, **Posto** can serve as a valuable resource for researchers and
17 engineers working on the design of modern autonomous systems. By integrat-
18 ing **Posto** into their current processes, they can evaluate the safety of their
19 designs. They can then analyze the resulting safety guarantees or violations
20 (counterexamples) to gain deeper insights and improve their designs.

21 **5. Conclusions and Future Work**

22 We introduced here **Posto**, an open-source Python-based software, aim-
23 ing at performing statistical monitoring for safety verification of complex au-
24 tonomous systems, including those driven by deep neural networks (DNN).
25 While traditional monitoring tools depend on formal models to assess safety,
26 **Posto** only requires the system’s input–output behavior (I/O execution) and
27 the log (which may be noisy and incomplete). Our experiments demonstrate
28 that **Posto** can efficiently monitor reasonably complex systems, including
29 those equipped with DNN-based controllers. We believe **Posto** will be val-
30 uable not only for engineers to analyze system logs to detect safety violations—
31 such as collision risks or other undesirable behaviors in robotics and au-
32 tonomous vehicles—but also for researchers developing new monitoring-based
33 methods. In such cases, **Posto** can serve as a useful comparative baseline.

34 Given the scalability of our method and the observations it offer, our
35 implementation paves the way for several future research directions. Our
36 ultimate aim is to adopt a fully representation-free strategy, where the I/O
37 execution representation will be co-designed with the monitoring approach.

1 **References**

- 2 [1] Free Software Foundation, Gnu general public license, version 3.
3 URL <https://www.gnu.org/licenses/gpl-3.0.en.html>
- 4 [2] A. Lab, Posto: Probabilistic Safety Monitoring – User Manual, Techni-
5 cal report, University of Alabama, User Manual (2024).
6 URL <https://autmn-lab.github.io/Posto/files/Posto-manual.pdf>
- 8 [3] B. Ghosh, É. André, MoULDyS: Monitoring of autonomous systems in
9 the presence of uncertainties, Science of Computer Programming 230
10 (Aug. 2023). [doi:10.1016/j.scico.2023.102976](https://doi.org/10.1016/j.scico.2023.102976).
- 11 [4] B. Ghosh, É. André, Probabilistic safety verification of distributed sys-
12 tems: A statistical approach for monitoring, in: C. Ferreira, C. A.
13 Mezzina (Eds.), FORTE, Vol. 15732 of Lecture Notes in Computer Sci-
14 ence, Springer, 2025, pp. 114–133. [doi:10.1007/978-3-031-95497-9_7](https://doi.org/10.1007/978-3-031-95497-9_7).
- 16 [5] D. A. Basin, G. Caronni, S. Ereth, M. Harvan, F. Klaedtke, H. Man-
17 tel, Scalable offline monitoring of temporal specifications, Formal
18 Methods in System Design 49 (1-2) (2016) 75–108. [doi:10.1007/s10703-016-0242-y](https://doi.org/10.1007/s10703-016-0242-y).
- 20 [6] E. Bartocci, J. V. Deshmukh, A. Donzé, G. E. Fainekos, O. Maler,
21 D. Ničković, S. Sankaranarayanan, Specification-based monitoring of
22 cyber-physical systems: A survey on theory, tools and applications,
23 in: E. Bartocci, Y. Falcone (Eds.), Lectures on Runtime Verification
24 – Introductory and Advanced Topics, Vol. 10457 of Lecture Notes in
25 Computer Science, Springer, 2018, pp. 135–175. [doi:10.1007/978-3-319-75632-5_5](https://doi.org/10.1007/978-3-319-75632-5_5).
- 27 [7] B. Ghosh, É. André, Offline and online monitoring of scattered uncer-
28 tain logs using uncertain linear dynamical systems, in: M. Mousavi,
29 A. Philippou (Eds.), FORTE, Vol. 13273 of Lecture Notes in Computer
30 Science, Springer, 2022, pp. 67–87. [doi:10.1007/978-3-031-08679-3_5](https://doi.org/10.1007/978-3-031-08679-3_5).
- 32 [8] B. Ghosh, É. André, Offline and online energy-efficient monitoring of
33 scattered uncertain logs using a bounding model, Logical Methods in
34 Computer Science 20 (1) (2023) 2:1–2:33. [doi:10.46298/lmcs-20\(1:2\)2024](https://doi.org/10.46298/lmcs-20(1:2)2024).

- 1 [9] M. Waga, T. Akazaki, I. Hasuo, A Boyer-Moore type algorithm for timed
 2 pattern matching, in: M. Fränzle, N. Markey (Eds.), FORMATS, Vol.
 3 9884 of Lecture Notes in Computer Science, Springer, 2016, pp. 121–139.
 4 [doi:10.1007/978-3-319-44878-7_8](https://doi.org/10.1007/978-3-319-44878-7_8).
- 5 [10] A. Bakhirkin, T. Ferrère, D. Ničković, O. Maler, E. Asarin, Online timed
 6 pattern matching using automata, in: D. N. Jansen, P. Pavithra (Eds.),
 7 FORMATS, Vol. 11022 of Lecture Notes in Computer Science, Springer,
 8 2018, pp. 215–232. [doi:10.1007/978-3-030-00151-3_13](https://doi.org/10.1007/978-3-030-00151-3_13).
- 9 [11] M. Waga, É. André, I. Hasuo, Parametric timed pattern matching, ACM
 10 Transactions on Software Engineering and Methodology 32 (1) (2023)
 11 10:1–10:35. [doi:10.1145/3517194](https://doi.org/10.1145/3517194).
- 12 [12] M. Waga, É. André, I. Hasuo, Model-bounded monitoring of hybrid
 13 systems, ACM Transactions on Cyber-Physical Systems 6 (4) (2022)
 14 30:1–30:26. [doi:10.1145/3529095](https://doi.org/10.1145/3529095).
- 15 [13] N. Halbwachs, Y.-É. Proy, P. Raymond, Verification of linear hybrid
 16 systems by means of convex approximations, in: B. Le Charlier (Ed.),
 17 SAS, Vol. 864 of Lecture Notes in Computer Science, Springer, 1994,
 18 pp. 223–237. [doi:10.1007/3-540-58485-4_43](https://doi.org/10.1007/3-540-58485-4_43).
- 19 [14] A. Legay, B. Delahaye, S. Bensalem, Statistical model checking: An
 20 overview, in: H. Barringer, Y. Falcone, B. Finkbeiner, K. Havelund,
 21 I. Lee, G. J. Pace, G. Rosu, O. Sokolsky, N. Tillmann (Eds.), RV, Vol.
 22 6418 of Lecture Notes in Computer Science, Springer, 2010, pp. 122–135.
 23 [doi:10.1007/978-3-642-16612-9_11](https://doi.org/10.1007/978-3-642-16612-9_11).
- 24 [15] A. Lab, Posto: Probabilistic Safety Monitoring – Recreating Results,
 25 Technical report, University of Alabama, Recreating Results (2024).
 26 URL [https://autmn-lab.github.io/Posto/files/
 27 Posto-recreation.pdf](https://autmn-lab.github.io/Posto/files/Posto-recreation.pdf)
- 28 [16] R. Ivanov, J. Weimer, R. Alur, G. J. Pappas, I. Lee, Verisig: Verifying
 29 safety properties of hybrid systems with neural network controllers, in:
 30 N. Ozay, P. Prabhakar (Eds.), HSCC, ACM, 2019, pp. 169–178. [doi:
 31 10.1145/3302504.3311806](https://doi.org/10.1145/3302504.3311806).




Figure A.1: **Posto Architecture and Dataflow.** Each block represents a core architectural component. Inputs are shown as hollow rectangles on the left, internal components as solid gray rectangles, and outputs as purple arrows on the right. Interactions between components are indicated by solid blue arrows.

1 Appendix A. Recreating Existing Results

2 This appendix lists the exact commands required to recreate the figures
3 reported in the paper. All commands are executed from the root directory of
4 the **Posto** repository using the `artEval.py` script ³. Due to the stochastic
5 nature of the proposed monitoring approach, regenerated figures are stochas-
6 tically equivalent but not necessarily identical.

7 *Environment Setup.*

```
8 cd /path/to/Posto
9 export POSTO_ROOT_DIR=/path/to/Posto
```

10 *Figures A.2(a)-(c).* To generate Fig. A.2, run:

³https://github.com/autmn-lab/posto_tool/blob/master/artEval.py




Figure A.2: **Jet Model Case Study** [4, Fig. 2]. *Random Trajectories* (Fig. A.2a). x - and y - axis represents state variables, and z -axis represents time step. *Impact of Logging Probability on Generation of Valid Trajectories* (Fig. A.2b). We illustrate how varying the logging probability affects the monitoring execution time and the percentage of valid trajectories out of the total number of generated trajectories. *Impact of Probabilistic confidence c* (Fig. A.2c). We illustrate the impact of logging probability on the time taken to perform monitoring and safety inferences. The red dot denotes an instance where the system was inferred to be unsafe with a counterexample, and green dots denote instances where the system was inferred to be probabilistically safe with the corresponding confidence value indicated in the x -axis.

```

1 python artEval.py --fig=A2a
2 python artEval.py --fig=A2b
3 python artEval.py --fig=A2c

```

4 *Figures A.3(a)-(d).* To generate Fig. A.3, run one of the following commands:

```

5 python artEval.py --fig=A3a
6 python artEval.py --fig=A3b
7 python artEval.py --fig=A3c
8 python artEval.py --fig=A3d

```

9 *Figures A.4(a)-(d).* To generate Fig. A.4, run one of the following commands:

```

10 python artEval.py --fig=A4a
11 python artEval.py --fig=A4b
12 python artEval.py --fig=A4c
13 python artEval.py --fig=A4d

```

14 Note that for each command, the auxiliary plot is displayed first, followed by
15 the corresponding Fig. A.4 subfigure. Closing the first window reveals the
16 target figure.

17 *Figure B.5.* To generate Fig. B.5, run:

```
18 python artEval.py --fig=B5
```




Figure A.3: Jet Monitoring [4, Fig. 3]. Evolution of the state— x for the jet model—with time. The volume of the records increases from left to right, and the probability of logging increases from bottom to top. The colored trajectories are the generated random valid trajectories during the process of monitoring, the dotted trajectory in red is an unsafe trajectory discovered during the process of monitoring (see Fig. A.3d), the black regions are records in the log given as an input to the monitoring algorithm, and the dark-brown regions are unsafe-records in the log (Fig. A.3b). The red dotted line represents the safe range for the variable.

1 Appendix B. Mountain Car Case Study

2 In this appendix, we report new results for a Mountain Car Controller
 3 that uses a neural-network-based policy [16]. The controller, provided as
 4 a trained .h5 model, is executed using the ANN mode of **Posto** and maps
 5 the car’s position and velocity to an acceleration command which is used to
 6 compute the next position and velocity of the car [16]. The discretized closed-
 7 loop dynamics are subject to injected environmental uncertainty. Safety is
 8 defined by requiring the velocity to remain below $v \leq 0.055$ at all time steps.
 9 Logs are generated under four configurations combining logging probabilities
 10 of 20% and 40% with uncertainty levels of 0.004 and 0.008. Only the con-
 11 figuration with 40% logging probability and uncertainty 0.008 violates the
 12 safety constraint, while the remaining cases are inferred to be safe, as shown
 13 in Fig. B.6.

14 Recreating Figure B.6

15 Fig. B.6 reports the monitoring results for the Mountain Car controller
 16 across all four logging and uncertainty configurations. To recreate all subfig-




Figure A.4: **Van der Pol Oscillator Monitoring** [4, Fig. 4]. Evolution of the state— y for the Van der Pol model—with time. The volume of the records increases from left to right, and the probability of logging increases from bottom to top. The colored trajectories are the generated random valid trajectories during the process of monitoring, the dotted trajectory in red is an unsafe trajectory discovered during the process of monitoring (see Fig. A.4d), the black regions are records in the log given as an input to the monitoring algorithm, and the dark-brown regions are unsafe-records in the log (Fig. A.4b). The red dotted line represents the safe range for the variable.

1 ures of Fig. B.6, run the `artEvalNN.py` script ⁴ using the following commands
2 from the root directory of the **Posto** repository:

```

3   python artEvalNN.py --fig=B6a
4   python artEvalNN.py --fig=B6b
5   python artEvalNN.py --fig=B6c
6   python artEvalNN.py --fig=B6d

```

7 Each command generates the corresponding subfigure of Fig. B.6, reflecting
8 one combination of logging probability and uncertainty level. Due to the
9 stochastic nature of the monitoring procedure, regenerated plots may dif-
10 fer slightly across runs while preserving the qualitative behavior and safety
11 outcomes reported in the paper.

⁴https://github.com/autmn-lab/posto_tool/blob/master/artEvalNN.py




Figure B.5: **Log Generation Example.** The reference trajectory is shown in blue, and the brown boxes illustrate the corresponding noisy samples of the log obtained using given logging probability and noise level.




Figure B.6: Mountain Car Controller Monitoring. Evolution of the state— v for the Mountain Car Controller—with time. The volume of the records increases from left to right, and the probability of logging increases from bottom to top. The colored trajectories are the generated random valid trajectories during the process of monitoring, the black regions are records in the log given as an input to the monitoring algorithm, and the dark-brown regions are unsafe-records in the log (Fig. B.6b). The red dotted line represents the safe range for the variable.